Bayesian Network model of NHTSA: FARS Data

Joseph Miller

Rutgers University arch1190@gmail.com

June 2, 2016



Overview

Introduction

FARS Overview

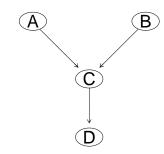
Preparing the Data

Building the Model



Introduction to Belief Networks

- Encodes random variables and their conditional (in)dependencies. In general, Pr(X|par(X))
- Can dramatically reduce the number of parameters (example: 15 vs. 8 if binary data) and model size.
- All BNs are DAGs.
- Different DAGs can correspond to the same factorization.



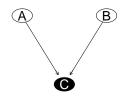
$$Pr(A, B, C, D) = Pr(D|C) Pr(C|A, B) Pr(A) Pr(B)$$



Introduction



In a causal chain, when B is known, C is independent of A: information "stops" at B. $Pr(C|B) Pr(B|A) Pr(A) \Rightarrow$ Pr(C, A|B) = $Pr(C|B) Pr(A) \times \left[\frac{Pr(B|A)}{Pr(B)}\right]$



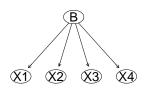
When C is known, A and B may be dependent: information "flows" through C. $Pr(A) Pr(B) Pr(C|A, B) \Rightarrow$ Pr(A, B|C) = $Pr(B) Pr(A|B) \times \left[\frac{Pr(C|B)}{Pr(C)}\right]$

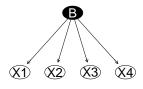
Coin Flip Model

Introduction

- ► Consider an IID coin flip experiment with $Pr(H) = \theta$.
- B is a latent variable that we can model hierarchically where X_i ~ Bern(θ).
- ▶ Is X_4 independent of the event $(X_1, X_2, X_3) = (H,H,H)$?
- ▶ What about if $\theta = 1/2$?

BNs generalize hierarchical modeling

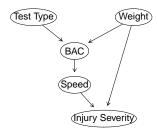




BN Sketch

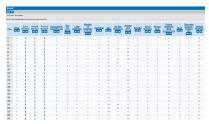
Introduction

- Plausible hierarchical structure in the data.
- BAC only acts on Injury Severity through Speed and correlates with it through Weight.
- Number of passengers (not shown) influences number injured.
- Passengers influence driver through distraction (not shown).



oduction FARS Overview Preparing the Data Building the Model

BN Sketch



Goal:

Build a model to predict VFATCOUNT (vehicle fatality count).

- 1. Prepare Data
- 2. Structure Learning
- 3. Parameter Learning
- 4. Evaluate Model

Difficulties:

- Mixed categorical and continuous
- Categorical data with many levels and few observations in some levels
- Many different missing value codes
- > 100 variables in the entire dataset
- Observations with multiple codes



Profiling the Data

> num_categories					
statenum	vnumber	atmcond	atmcond2	accmon	dayofweek
51	40	13	13	12	7
fhevent	lightcond	numfatal	rfun	speeding	age
51	9	8	16	2	102
airbag	alcres	alcsts	alctst	restraintmisuse	drugres1
12	68	6	13	3	99
drugres2	drugres3	drugtst1	drugtst2	drugtst3	ejection
66	57	- 9	- 8	9	8
injury	methalc	methdrug	ptype	druginv	alcinvol
8	7	6	3	- 5	4
гасе	restraint	seatpos	sex	inimpact	mhevent
20	17	12	5	25	51
vfatcount	numoccs	rollover	travspd	laccdate	malcohol
8	42	5	140	51	2
feet	inches	prevdwi	prevoth	spdrel	crashtype
8	14	10	16	7	88
dridistract	drivisobs	spdlim	dridistract2	dridistract3	drivisobs2
24	20	20	11	2	11
drivisobs3					



Vehicle Fatality Count

1 2 3 4 5 6 7	1 2 3 4 5	19903 23083 1424 200 54 20	
6			
7 8	6 9	1 1	

- ► Variable: VFATCOUNT
- High entropy
- Will bin higher values

Cleaning Data

```
strings <- strsplit(as.character(fars$drivisobs),
split = " ")

l <- vector(mode="integer", length=length(strings))
for (i in 1:length(strings)) { # Get vector with
lengths of strings.
    l[i] <- length(strings[[i]])
}
table(l) # Histogram of number of values in each entry</pre>
```

```
l
1 2 3
44650 30 6
```

Transformations

New Variables

- Four new columns, two each from DRIVISOBS and DRIDISTRACT.
- ► Transformed INCHES (0-12) and feet (2-7) to continuous variable HEIGHT.

Remaining Work

- 1. Eliminate highly unbalanced features.
- 2. Merging categories (will handle missing values).
- 3. Structure/Parameter Learning.
- 4. Evaluate Models.



Low Variability Features

Heading

- ▶ 13/54 variables with near-zero-variability:
 - < 10% unique values.</p>
 - Ratio of frequency of most common to second most common category
 20:1.
- Parameter estimates too variables.
- Will handle metric variables differently.



Merging Categories

```
arrange(drugRes,
Freq)[1:10,]
```

```
Drug Freq
    111
    115
3
    117
4
    143
5
    182
6
    196
    198
8
    200
9
    204
10
    229
```

```
Drug Freq
11
    233
12
    248
13
    255
    336
14
15
    338
    342
16
17
    349
18
    357
19
    386
20
    396
21
    403
22
    408
23
    409
24
    504
```



Merging Categories

- Drug Test Results has 99 observed categories, 150+ possible.
 Others bad, too.
- Want:
 - Combine levels of cat. variables that are most similar in relation to target variable.
 - Depenent on learned structure of BN, which is dependent on how the categories are merged.
 - Data is not ordinal (merge categories).
 - Reminiscent of cluster analysis.
- Am not aware of an appropriate/efficient algorithm.
- Forfeit my ability to predict elements within these categories.



Merging Categories Below Cutoff Freq

```
merge_categories <- function(fdata, p) {
   merged_levs <- list()
   for (col in 1:dim(fdata)[2]) {
     dat <- fdata[,col]
     lowfreq <-
names(which(prop.table(table(dat)) < p))
     levels(dat)[levels(dat) %in% lowfreq] <-
"Other"
     merged_levs[[names(fdata)[col]]] <- lowfreq
     fdata[,col] <- dat
   }
   return(list(fdata, merged_levs))
}</pre>
```

Merging Categories Below Cutoff Freq

```
out <-
merge_categories(fdata,
0.01)
fdata <- out[[1]];
merged_levs <- out[[2]]</pre>
```

- ► Smallest allowed factor is now 0.01 × 44686.
- Much room for improvement exists!

```
arrange(drugRes,
Freq) [1:10,]
    Drug
           Freq
     600
            471
2
3
     999
            860
     605
            879
     996
            974
      95
           1036
     997
           1237
           1426
8
   0ther
           4065
9
           9953
10
        0 23785
```

Discretizing Metric Data

```
mdata <- discretize(data
= temp, method =
"hartemink", breaks = 7,
ibreaks = 30,
idisc="quantile")</pre>
```

- Cuts numerical data into 30 quantiles, combines intervals in a way that reduces pairwise mutual information minimally.
- Final intervals depend on other variables in data frame: made 6-variable drame with 5 metric variables and the target variable.

levels(mdata\$height)

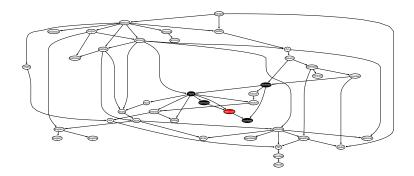
```
[1] "(33.9998,65.0003]"
"(65.0003,68.9999]"
"(68.9999,75.0007]"
[4] "(75.0007,83.9991]"
"(83.9991,84.001]"
```

Structure Learning

- ► Efficient algorithms exist to learn tree structures. Arbitrary structures: search space is too big.
- Two notable methods for comparing structures:
 - Independence tests: reduce search space by restricting to learning trees.
 - Scoring criterion: can more efficiently search structures. No guarantee of convergence. Must choose search algorithm.
- Chose BIC/MDL and Bayesian Dirichlet Method. Space of DAGs must still be searched.
 - Greedy Search with perturbations and restarts.

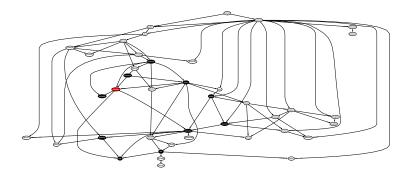


BIC Optimized Model





Bayesian Model





Parameter Learning and Classification

- Parameter learning is straightforward: MAP estimates with uniform Dirichlet prior on all parameters.
- Start with predictions given Markov Blanket (black nodes) on target variable: number of fatalities in driver's vehicle.
- ▶ BIC Model: 185/200 correct classifications, 42 seconds.
- ▶ Bayesian Model: 194/200 correct classifications, 191 seconds.
- ▶ 200 samples, low SE.

> table(test[sample200,]\$vfatcount)

```
0 1 2 Other
94 96 8 2
```



troduction FARS Overview Preparing the Data **Building the Model**

Summary: BIC-scored Model

Confusion Matrix and Statistics

Actual Predicted 0 1 2 Other 0 82 2 0 0 1 11 93 0 0 2 1 1 8 0 Other 0 0 0 2

Overall Statistics

```
Accuracy : 0.925
95% CI : (0.8793, 0.9574)
No Information Rate : 0.48
```

P-Value [Acc > NIR] : < 2.2e-16

Kappa : 0.8639 Mcnemar's Test P-Value : NA

Statistics by Class:

	Class: 0	Class: 1	Class: 2	Class: Other
Sensitivity	0.8723	0.9688	1.0000	1.00
Specificity	0.9811	0.8942	0.9896	1.00
Pos Pred Value	0.9762	0.8942	0.8000	1.00
Neg Pred Value	0.8966	0.9688	1.0000	1.00
Prevalence	0.4700	0.4800	0.0400	0.01
Detection Rate	0.4100	0.4650	0.0400	0.01
Detection Prevalence	0.4200	0.5200	0.0500	0.01
Balanced Accuracy	0.9267	0.9315	0.9948	1.00



troduction FARS Overview Preparing the Data **Building the Model**

Summary: Bayesian Model

Confusion Matrix and Statistics

```
Actual
Predicted 0 1 2 0ther
0 90 1 0 0
1 4 94 0 0
2 0 1 8 0
Other 0 0 0 2 2
```

Overall Statistics

```
Accuracy: 0.97
95% CI: (0.9358, 0.9889)
No Information Rate: 0.48
P-Value [Acc > NIR]: < 2.2e-16
```

Kappa : 0.9454 Mcnemar's Test P-Value : NA

Statistics by Class:

```
Class: 0 Class: 1 Class: 2 Class: Other
Sensitivity
                      0.9574
                                0.9792
                                         1.0000
                                                        1.00
Specificity
                                0.9615
                      0.9906
                                         0.9948
                                                        1.00
Pos Pred Value
                      0.9890
                                0.9592
                                         0.8889
                                                        1.00
                                        1.0000
Neg Pred Value
                      0.9633
                                0.9804
                                                        1.00
Prevalence
                      0.4700
                                0.4800
                                         0.0400
                                                        0.01
Detection Rate
                      0.4500
                                0.4700
                                         0.0400
                                                        0.01
Detection Prevalence 0.4550
                                0.4900
                                         0.0450
                                                        0.01
Balanced Accuracy
                                0.9704
                                         0.9974
                                                        1.00
                       0.9740
```



Evaluating Forecasts

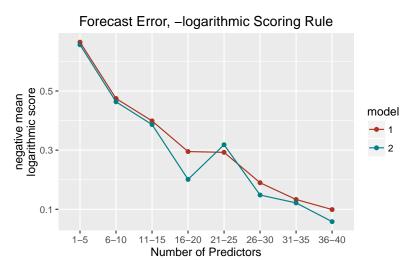
Algorithm:

- 1. *N*, picked uniformly from 1 to 40 (41 total variables).
- Choose N PREDICTORS, picked without replacement from the columns.
- 3. Predict VFATCOUNT with each model.
- 4. Repeated (1-3) 1000 times.
- Calculate average log(p) of predicted class, a proper scoring rule.

Results:

- Correct classification rate:
 - **▶** 85.8%
 - **▶** 86.2%
- ▶ average log(p)
 - ► -0.3226594
 - -0.2969937

Error Trend





References

- bnlearn R package and documentation, Marco Scutari
- gRain R package and documentation, Soren Hojsgaard
- ggplot2 R package, Hadley Wickham
- dplyr R package, Hadley Wickham



The End